

Deep Learning

Exercise 6: MNIST and LeNet in PyTorch

Instructor: Manuel Günther

Email: guenther@ifi.uzh.ch

Office: AND 2.54

Friday, April 16, 2020

Outline

- 1 PyTorch
- 2 MNIST Training with PyTorch

Outline

- 1 PyTorch
 - Installation
 - Building a Network
 - Datasets and Batches
 - The Training Loop
 - Running on the GPU

Installation

Cuda Installation

- Install Cuda driver (if you have a Cuda-enabled GPU)
 - List of supported GPUs: <https://developer.nvidia.com/cuda-legacy-gpus>
 - Install driver for your OS <https://www.nvidia.com/Download>

PyTorch Installation

- Install from `pytorch` channel via Conda:
`conda install pytorch torchvision cudatoolkit -c pytorch`
- Extensive documentation and tutorials:
 - Tutorials: <https://pytorch.org/tutorials>
 - Reference documentation: <https://pytorch.org/docs/stable>

Building a Network

Types of Layers

- Fully connected layer
`torch.nn.Linear`
 - `in_features = D`
 - `out_features = K`
 - `bias = True`
- Convolutional layer
`torch.nn.Conv2d`
 - `in_channels = C`
 - `out_channels = Q`
 - `kernel_size = U ∨ (U, V)`
 - `padding, stride, bias`

Related Functions

- Activation functions:
`torch.nn.Sigmoid`,
`torch.nn.Tanh`,
`torch.nn.Softmax` ...
- Pooling:
`torch.nn.MaxPool2d`,
`torch.nn.AvgPool2d`
 - `kernel_size, padding`
- Input flattening:
`torch.nn.Flatten`

Building a Network

Defining Network Topology

- Derive class from `torch.nn.Module`
- Instantiating layers and functions in constructor `__init__(self)`
 - Separate instances for each layer
 - One instance for each function

Executing Network

- Implement `forward(self, x)`
 - Call your layers sequentially
 - Return last output
- `backward` not required
 - Details in next lecture

Example: Non-Linear Regression

```
class Regression (torch.nn.Module):  
    def __init__(self, K):  
        # call base class constructor  
        super(Regression,self).__init__()  
  
        # initialize components  
        self.W1 = torch.nn.Linear(1, K, bias=True)  
        self.activation = torch.nn.Sigmoid()  
        self.w2 = torch.nn.Linear(K, 1)  
  
    def forward(self, x):  
        a = self.W1(x)  
        h = self.activation(a)  
        z = self.w2(h)  
        return z  
  
    def forward(self, x):  
        return self.w2(self.activation(self.W1(x)))
```

Datasets and Batches

Datasets

- Many available in `torchvision.datasets`
 - `torchvision.datasets.MNIST`
 - `torchvision.datasets.ImageNet`
- Common interface:
 - `root`: Directory of raw data
 - `train`: set to `False` for test set
 - `download`: downloads data if required
 - `transform`: preprocessing
- Return `PIL` images

Transforms

- Prepares the input
 - Depends on original data
- Implemented in `torchvision.transforms`
 - `Resize((D, E))`: height, width
 - `Normalize(mean, std)`
 - `Lambda(callable)`: generic
 - `ToTensor`: `PIL` → `tensor`
- Combining several transforms:
 - `Compose((trans1,trans2))`
- Additionally: target transforms

Datasets and Batches

Data Loader

- `torch.utils.data.DataLoader`
 - `dataset`: see above
 - `batch_size` = B
 - `shuffle` after each epoch
 - `num_workers`: parallel execution on the CPU

Example MNIST Dataset

```
# obtain datasets
transform = torchvision.transforms.ToTensor()
train_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=True, download=True,
    transform=transform
)
test_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=False, download=True,
    transform=transform
)

# loaders
train_loader = torch.utils.data.DataLoader(
    train_set, shuffle=True, batch_size=64
)
test_loader = torch.utils.data.DataLoader(
    test_set, shuffle=False, batch_size=100
)
```


The Training Loop

Loss Functions

- `torch.nn.MSELoss`: \mathcal{J}^{L_2}
- `torch.nn.BCEWithLogitsLoss`:
binary \mathcal{J}^{CE} and σ activation
- `torch.nn.CrossEntropyLoss`:
categorical \mathcal{J}^{CE} and SoftMax

Learning Strategy

- `torch.optim.SGD`
 - `params`: All optimizable weights
 - `lr`: Learning rate η
 - `momentum`: μ

Example Training Loop

```
# instantiate everything
network = Network()
loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(
    params=network.parameters(),
    lr=0.01, momentum=0.9
)

# training epoch
for epoch in range(epochs):
    for x,t in train_loader:
        # DO NOT FORGET:
        optimizer.zero_grad()
        z = network(x)
        J = loss(z, t)
        J.backward()
        optimizer.step()
    # compute test accuracy
```

Running on the GPU

Preparation

- Test cuda availability:
`torch.cuda.is_available()`
- Select device "cpu" or "cuda":
`device = torch.device("cuda")`
- Move everything to the device:
`network.to(device)`
`x.to(device)`
`t.to(device)`

Speed Warning

Can be slow on CPU (1 min per epoch)

Example Training Loop

```
# instantiate everything
device = torch.device("cuda")
network = Network().to(device)
loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(
    params=network.parameters(),
    lr=0.01, momentum=0.9
)

# training epoch
for epoch in range(epochs):
    for x,t in train_loader:
        optimizer.zero_grad()
        z = network(x.to(device))
        J = loss(z, t.to(device))
        J.backward()
        optimizer.step()
    # compute validation accuracy
```

Outline

2 MNIST Training with PyTorch

MNIST Training with PyTorch

Task 1

- ❶ Implement a network in PyTorch including:
 - One fully-connected layer with 784 inputs and K hidden neurons
 - A sigmoid activation function
 - One fully-connected layer with K inputs and $O = 10$ logits
- ❷ Create training and test set for the MNIST dataset
 - Transform 28×28 image into vector of size 784
- ❸ Choose optimizer SGD with nicely-working parameters
- ❹ Train the network using categorical cross-entropy and SoftMax
- ❺ Compute and print test set accuracy after each epoch
- ❻ Train the network for 100 epochs and get highest test accuracy

MNIST Training with PyTorch

Task 2

- ① Improve fully-connected network with convolutions
 - ① Add one or more convolutional layers in the beginning
 - Select kernel parameters U, V , stride, padding
 - ② Add a pooling layer after each convolution layer
 - Select type of pooling and its parameters
 - ③ Have one fully-connected layer with $K = ???$ inputs and $O = 10$ logits
- ② Create a training and test set for the MNIST dataset
 - What kinds of transforms do you need here?
- ③ Choose optimizer SGD with nicely-working parameters
- ④ Train the network using categorical cross-entropy and SoftMax
- ⑤ Compare best test set accuracy to fully-connected model