

# Deep Learning

## Exercise 8: Open-Set Learning

Instructor: Manuel Günther

Email: [guenther@ifi.uzh.ch](mailto:guenther@ifi.uzh.ch)

Office: AND 2.54

Friday, April 30, 2020

# Outline

- 1 Open-Set Training
- 2 Evaluation
- 3 Autograd Implementation

# Outline

- 1 Open-Set Training
  - MNIST Split
  - Network Definition
  - Loss Functions

# MNIST Split

## MNIST Dataset

- MNIST total 10 classes
  - 10 different digits
- Split into 3 categories:
  - 4 known classes, e.g.:  
4, 5, 8, 9
  - 4 known unknown classes, e.g.:  
0, 2, 3, 7
  - 2 unknown unknown classes, e.g.:  
1, 6
- Train known and known unknown
- Evaluate all three categories

## Task 1: Target Vectors

- One-hot targets for knowns:
  - 4  $\Rightarrow (1, 0, 0, 0)$
  - 5  $\Rightarrow (0, 1, 0, 0)$
  - 8  $\Rightarrow (0, 0, 1, 0)$
  - 9  $\Rightarrow (0, 0, 0, 1)$
- Equal targets for unknowns:
  - 0  $\Rightarrow (.25, .25, .25, .25)$
  - 2  $\Rightarrow (.25, .25, .25, .25)$
  - 3  $\Rightarrow (.25, .25, .25, .25)$
  - 7  $\Rightarrow (.25, .25, .25, .25)$

# Datasets Loaders

## Example MNIST Dataset

```
# obtain datasets
transform = torchvision.transforms.ToTensor()
train_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=True, download=True,
    transform=transform
)
test_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=False, download=True,
    transform=transform
)
# loaders
train_loader = torch.utils.data.DataLoader(
    train_set, shuffle=True, batch_size=64
)
test_loader = torch.utils.data.DataLoader(
    test_set, shuffle=False, batch_size=100
)
```

## While Training

```
for inputs, targets in train_loader():
    # remove unknown unknown inputs
    ...
    # convert targets to one-hot
    targets = convert_target(targets)
    # forward inputs
    logits = network(inputs)
    # call loss function
    J = loss(logits, targets)
    ...
```

# Network Definition

## Convolutional Network

- Adaptation from Exercise 6
- 2 convolutions with padding
  - 1 input, 32 outputs,  $5 \times 5$  kernels
  - 32 input, 32 outputs,  $5 \times 5$  kernels
  - $2 \times 2$  Maximum pooling after each
  - Sigmoid activation after each
- Batch normalization (**important**)
- Linear layer with 50 outputs
- Sigmoid activation
- Linear layer with 4 outputs

## Task 2: Network Implementation

```
class Network():
    def __init__(self):
        self.conv1 = torch.nn.Conv2d(...)
        self.conv2 = torch.nn.Conv2d(...)
        self.pooling = torch.nn.MaxPool2d(...)
        self.activation = torch.nn.Sigmoid()
        self.bn = torch.nn.BatchNorm2d(...)

        self.fc1 = torch.nn.Linear(7*7*32, 50)
        self.fc2 = torch.nn.Linear(50, 4)

    def forward(self, input):
        ...
        return logits
```

# Loss Functions

## Default Implementation

- `torch.nn.CrossEntropyLoss` requires target class  $t$   
→ We have target vector  $\vec{t}$

## Easiest Implementation

$$\mathcal{J}^{\text{CE}} = - \sum_{o=1}^O t_o \ln y_o$$

## Other Implementation

$$\mathcal{J}^{\text{CE}} = - \sum_{o=1}^O t_o z_o + \frac{1}{O} \ln \sum_{o=1}^O e^{z_o}$$

## Task 3: Loss Function via Autograd

- To implement:  
`def loss(logits, targets):`
- $\ln \vec{y}$  from logits  
→ `torch.log_softmax`
- $\ln \sum_{o=1}^O e^{z_o}$   
→ `torch.logsumexp`
- + Remember that you get a batch  
→ Check the `dim` parameter
- To return: a `torch.tensor` with one element (check `torch.mean`)

# Outline

## 2 Evaluation



# Evaluation

## During Training

- Compute average confidence
  - $y_t$  for known samples
  - $1 - \max_o y_o + \frac{1}{O}$  for unknowns

## After Training

- Evaluate confidences  $\vec{y}$ 
  - False Positive Rate for unknown
  - Correct Classification Rate for known
- Use fixed threshold  $\tau = 0.5$

## False Positive Rate (FPR)

$$\frac{\left| \left\{ x^{[n]} \mid t^{[n]} = 0 \wedge \max_{1 \leq o \leq O} y_o^{[n]} \geq \tau \right\} \right|}{\left| \{ x^{[n]} \mid t^{[n]} = 0 \} \right|}$$

## Correct Classification Rate (CCR)

$$\frac{\left| \left\{ x^{[n]} \mid t^{[n]} > 0 \wedge \arg \max_{1 \leq o \leq O} y_o^{[n]} = t^{[n]} \wedge y_{t^{[n]}}^{[n]} \geq \tau \right\} \right|}{\left| \{ x^{[n]} \mid t^{[n]} > 0 \} \right|}$$

## Task 4

- Train network for 20 epochs
- Evaluate trained network

# Outline

- 3 Autograd Implementation
  - Autograd Extension in PyTorch

# Autograd Extension in PyTorch

## Loss and Derivative of $\mathcal{J}^{\text{CE}}$

$$\mathcal{J} = - \sum_{o=1}^O t_o \ln y_o \quad \frac{\partial \mathcal{J}}{\partial z_o} = y_o - t_o$$

## Implementation in PyTorch

- Implement a loss as `pytorch.autograd.Function`
- Online documentation
- Two static functions:
  - `forward(ctx, params)`
  - `backward(ctx, result)`
- Store elements in `ctx`

## Optional Task: Autograd Function

```
class CELoss(torch.autograd.function):
    @staticmethod
    def forward(ctx, logits, targets):
        # save for backward
        ctx.save_for_backward(logits, targets)
        # implement loss function
        J = ...
        return J

    @staticmethod
    def backward(ctx, J):
        # get stored values
        logits, targets = ctx.saved_tensors
        # compute derivatives
        gradient = ...
        # gradient for each input of forward
        return gradient, None

# assign loss function
loss = CELoss.apply
```